

Classifying Primary Emotions in Faces

Yuxuan Liu

mpanda@live.unc.edu

Evan Vitkus

vitkuse@live.unc.edu

Caleb Xu

cjxu@live.unc.edu

Matthew Zheng

mdz1999@live.unc.edu

ABSTRACT

The identification of face-like structures is a problem that had become well-researched in recent years. While facial-recognition software is becoming more and more accurate, the messages and implications contained in the emotions of the person are often lost when still images are captured. Additionally, this software is often lackluster when used on cartoon faces. In this paper, we create various models to attempt to analyze images, both real and cartoon, to identify the emotions portrayed in the faces based on a set of facial landmarks with models trained on real facial data reaching 62% and models trained on cartoon faces reach 98% accuracy. We were able to conclude that the way that emotions are portrayed in real faces are relatively consistent with other real faces and similar for cartoon faces. However, cartoon and real faces have little in common with their landmarks differing significantly with accuracy levels around 25%.

1. INTRODUCTION

As machine learning models have been on the forefront of research in nearly every field of study, the use cases and applications of such models has rapidly broadened. With the models able to form accurate predictions given pre-existing data, the use of such models is immensely attractive in fields that present difficult problems that have large amounts of existing data. One particular field of study that has attracted much interest in recent years is the identification of emotions. Understanding the emotions that a person is experiencing is critical in understanding the meaning and intention behind their disposition; however, as a naturally subjective trait, emotions are often misunderstood resulting in confusion and uncomfortable situations. While possible to detect a person's emotions based on their actions, this is often still subjective and is often too late.

Emotions are displayed through a variety of means known as affect displays. These affect displays are sometimes sufficient to determine emotion, but are not always comprehensive. Additionally, even if an emotion is perceived it is often difficult to understand why that perception occurs. While it may be obvious to say a person is expressing the emotion of joy when they are smiling, it doesn't capture every feature that the brain takes into account in processing the expression. In a different facial context, a smile could also indicate another emotion such as awkwardness or surprise.

With the subjective nature of emotions and the inability to properly identify the meaning of particular affect displays without additional information, it can be difficult to identify emotions that are artificially displayed. For example, when drawing, digital artists must render emotions as they believe they would appear. Then, the consumer will interpret the digitally produced face. This can result in an image that displays emotion incorrectly or in a manner that is difficult to perceive. As a result, messages can be warped or misinterpreted. It is therefore imperative that we create models capable of quantitatively interpreting emotions.

Current research has shown that an effective way of determining emotion is through the Facial Action Coding System (FACS). [7] This system makes use of action units, particular facial movements such as the raising of cheeks or the dropping of the jaw to identify the emotion expressed on a face. With animated or cartoon faces, however, these action units can not be relied on. These faces often lack certain of these action units and would therefore be unidentifiable with FACS.

Since we can not rely on FACS due to the lack of captured motion, we decided to analyze the known 68 facial landmarks. Using these landmarks, we can understand the position of pieces of the face that encompass the action units. Using 68 landmarks we seek to train models capable of identifying the emotions present in images of faces both real and cartoon. We then cross-test these models to see how closely real and cartoon faces' landmarks are to one another. Finally, we train a model on a combination of the data. Using these models, we seek to determine how closely related real and cartoon faces are, and how accurately emotions are displayed in them.

2. APPROACH

In order to detect emotions through facial expressions, we required many images of faces and a way to quantify their expressions. We acquired two datasets, the Facial Expression Research Group Database (FERG-DB) from Washington University and Karolinska Directed Emotional Faces (KDEF) from the Emotion Lab at the Karolinska Institutet. [1][5] The FERG-DB dataset is a collection of six cartoon people displaying a variety of emotions and the KDEF dataset is of 140 real people displaying the same set of emotions taken from a variety of angles. Both of these datasets' images had been previously classified with emotions. To quantify the facial landmarks in these, we created a histogram of oriented gradients (HOG).[2] With our

newly fashioned, quantified dataset we ran multiple classifier machine learning algorithms to produce sets of models to determine the accuracy of the emotions expressed. Using scikit-learn, we created five separate model sets where we trained and tested on cartoon to cartoon, real to real, cartoon to real, real to cartoon, and a combination of cartoon and real to a combination of cartoon and real respectively. [6]




Figure 1: Comparison of original face with marked face from KDEF dataset




Figure 2: Comparison of original and marked face from FERF-DB dataset

3. PROCEDURE

After obtaining image data from FERF-DB and KDEF datasets, we ensured that all of the images were the same size. The KDEF images feature subjects from frontal and side camera angles. To ensure consistency with angles of images used from the FERF-DB dataset, we filtered the KDEF images to use only images captured from a direct frontal angle.

The FERF-DB images were already at the resolution of 768×768 , while the KDEF images were at the resolution of 562×762 . To address this discrepancy, all of the images were padded with plain white on left and right to the resolution 762×762 before being rescaled to resolution 768×768 through bicubic resampling in order to match the FERF-DB image sizes. These image resizing operations are implemented in `image_resize.py`.

Using pre-trained model files for `dlib`, we then generated 68 landmark coordinates for each FERF-DB and KDEF image. [3][4] These landmark generation operations are implemented in `cartoon_bulk_face_detector.py` and `real_bulk_face_detector.py`, respectively.

These landmarks are stored in `real_landmarks.csv` and `cartoon_landmarks.csv`. The emotion labels for these images were also quantized from one to seven, corresponding to anger, disgust, fear, joy, neutral expression, sadness, and surprise, respectively, and these labels are then stored in `real_labels.csv` and `cartoon_labels.csv`. The order of entries in each landmarks file corresponds to that of the entries in the respective labels file.

These landmark and label sets are then processed in `csv_process.sh`. The cartoon and real landmarks and labels are shuffled, and then a new landmark and label file containing equal parts of cartoon and real data is generated. The cartoon and real landmarks and labels are then also divided by a 75-25 split into training and test data, respectively. Finally, we train machine learning models in `sklearn-template.py`. The script takes two arguments, each being one of the words `real`, `cartoon`, or `combined`. The first argument selects the dataset upon which to train the models, and the second argument selects the dataset upon which to test the models. For example, the invocation `python sklearn-template.py cartoon real` will train the models on the cartoon face dataset and test them on the real face dataset. Without any arguments, the script will train and test all possible combinations of real, cartoon, and combined datasets.

The Python script produces an output file, `sklearn_output.csv`, showing the percentage of accuracy of each of the models trained and tested in each combination. The script trains a number of models in `sklearn`, including a random forest ensemble classifier, Ada Boost classifier, bagging classifier, extra trees classifier, gradient boosting classifier, and decision tree classifier.

4. RESULTS

As seen in Table 1, our model performed stronger on the cartoon faces than the real faces for every model we ran. Comparing each model's performance, each model predicted emotion more accurately on the cartoon image than on the real face. Our justification is due to the variance among faces. Our cartoon data was a compilation of six individuals, each with over 9,000 images. The KDEF dataset involved 140 different participants with only 8 viable images each. Our strongest performing model was gradient boosting followed by random forest. We received the most accurate models when comparing among the same data. This follows because there is less variance when predicting among the same dataset. We also received a lower average prediction score from our combined datasets. The reason also has to do with additional variance between the real and cartoon faces. The combined model performs better than trying to directly cross-compare because we are training on the both datasets.

Table 1: Training results on datasets within themselves

Model	Real \rightarrow Real	Cartoon \rightarrow Cartoon
Random Forest	0.567346	0.981730
Ada Boost	0.432653	0.617876
Gradient Boosting	0.624489	0.956037
Decision Tree	0.424489	0.898318

Table 2: *Training results on datasets between each other*

Model	Real \rightarrow Cartoon	Cartoon \rightarrow Real
Random Forest	0.276512	0.208163
Ada Boost	0.262755	0.297959
Gradient Boosting	0.292452	0.277551
Decision Tree	0.230438	0.228571

Table 3: *Training results on combined dataset within itself*

Model	Combined \rightarrow Combined
Random Forest	0.542857
Ada Boost	0.318367
Gradient Boosting	0.585714
Decision Tree	0.402040

5. CONCLUSIONS

With the accuracy that we’ve obtained from the different models, we’ve established several things. To no surprise, real faces and cartoon faces proved to be solid datasets to use to classify themselves. Real to real achieved upwards of 62% accuracy while cartoon to cartoon achieved upwards of 98% accuracy in classification with certain models. While real to real produces a weaker classifier, we believe that is due to the small size of our real face dataset.

The part that is interesting is that classifying one type using the other type as a training set is not very accurate. Most models using one set to classify the other barely does better than random guessing. Even trying to combine the training sets produced lackluster, although better, accuracy. The improved accuracy, however, can be attributed to the large size of the FERF-DB data set allowing for a large number of cartoon faces to be classified correctly.

One possible explanation for the inefficacy of the cross-tested models could be that we simply didn’t have enough data in the KDEF data set. This explanation doesn’t hold up well, however, because we had a decently large training set when using the cartoon set to train but it still produced an inaccurate classification for real datasets.

Another possible explanation is the radical difference between the two datasets. As a result, the model might as well be guessing as the landmarks on cartoon faces potentially do not match up or align at all with their real counterparts. In fact, visual inspection shows that the landmark features on the cartoon faces are often exaggerated compared to their realistic counterpart.

We may be tempted to conclude that there actually exists some amount of correlation between the landmarks on the real faces and on the cartoon faces. However, the accuracy of the cross-tested models is not significant enough to account for the randomization of the classification that inherently existed in the datasets. Since people are inherently imperfect at identifying and/or creating human facial expressions, it is unsafe to assume that the results are significant. The evidence only supports the claim that cartoon faces bear little resemblance to their real counterparts.

6. FUTURE WORK

Possible follow up to improve our models would be increasing the number of facial landmarks and a general increase in the size of the real face image data set. A larger dataset

for real faces would grant our model additional data to train with, resulting in better classification power. Additionally, with a larger number of facial landmarks per image, additional information about each image would be added as quantifiable information. More landmarks would result in a better depiction of the face by providing our data with additional contours to each face.

Another interesting way to interpret these findings would be to create or find a dataset where each real face has a cartoon face counterpart that is produced to be a replica of the real face. This will not only produce a more accurate model but will also better answer the question of whether people could intentionally produce artificial emotions in cartoons.

Additional interesting directions to continue with our research include adding images of faces in a natural environment, which includes having faces that are not oriented directly towards the camera and the generation of more natural cartoon faces displaying particular emotions.

7. ACKNOWLEDGEMENTS AND CITATIONS

We would like to acknowledge Jorge Silva for instructing us in COMP 562 and providing with the starting ground for this research. We also thank Washington University and the Karolinska Institutet for allowing us to use their data.

8. REFERENCES

- [1] D. Aneja, A. Colburn, G. Faigin, L. Shapiro, and B. Mones. Modeling stylized character expressions via deep learning. In *Asian Conference on Computer Vision*, pages 136–153. Springer, 2016.
- [2] A. Geitgey. Machine learning is fun! part 4: Modern face recognition with deep learning, Jul 2016.
- [3] D. King. dlib-models. Github repository for pretrained dlib-models.
- [4] D. E. King. Dlib-ml: A machine learning toolkit. *Journal of Machine Learning Research*, 10:1755–1758, 2009.
- [5] F. A. . Ö. A. Lundqvist, D. The karolinska directed emotional faces - kdef. 1998. real faces data retrieved here. Lundqvist, D., Flykt, A., & Öhman, A. (1998). The Karolinska Directed Emotional Faces - KDEF, CD ROM from Department of Clinical Neuroscience, Psychology section, Karolinska Institutet, ISBN 91-630-7164-9.
- [6] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [7] Y.-I. Tian, T. Kanade, and J. Cohn. Recognizing action units for facial expression analysis. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 23(2):97–115, Feb 2001.